

Answer 1

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn import mixture
from sklearn.model_selection import KFold
from scipy.stats import multivariate_normal
```

```
In [2]: def compute_log_likelihood(m,data):
    '''
    m is the model
    data is samples
    '''
    return m.score(data)
```

```

In [3]: def generate_GMM_samples(prior,number_of_samples,sig1,sig2,sig3,sig4,u1,
    u2,u3,u4):
    '''
    Args:
    prior of class 1 = prior[0]
    prior of class 2 = prior[1]
    prior of class 3 = prior[2]
    prior of class 4 = 1-prior[0]-prior[1]-prior[2]

    number_of_samples

    class 1- u_1, sig_1
    class 2- u_2, sig_2
    class 3- u_3, sig_3
    class 4- u_4, sig_4

    x is samples from zero-mean identity-covariance Gaussian sample generators

    generating class 1- A1*x+b1
    generating class 2- A2*x+b2
    generating class 3- A3*x+b3
    generating class 4- A4*x+b4

    '''
    from matplotlib.pyplot import figure
    #txt="Plot of data sampled from 4 gaussians "
    fig = plt.figure(figsize=(10,10));
    #fig.text(.35,0.05,txt,fontsize=15);

    samples_class1=[]
    samples_class2=[]
    samples_class3=[]
    samples_class4=[]

    sig_1=np.matrix(sig1)
    sig_2=np.matrix(sig2)
    sig_3=np.matrix(sig3)
    sig_4=np.matrix(sig4)

    u_1=np.matrix(u1).transpose()
    u_2=np.matrix(u2).transpose()
    u_3=np.matrix(u3).transpose()
    u_4=np.matrix(u4).transpose()

    prior=prior
    A1=np.linalg.cholesky(sig_1)
    b1=u_1

    A2=np.linalg.cholesky(sig_2)
    b2=u_2

    A3=np.linalg.cholesky(sig_3)
    b3=u_3

    A4=np.linalg.cholesky(sig_4)

```

```

b4=u_4

zero_mean=[0,0]
cov=[[1,0],[0,1]]

for i in range(number_of_samples):
    uniform_sample=np.random.uniform()

    sample_from_zero_mean_identity_covariance=np.random.multivariate
_normal(zero_mean,cov,[1]).transpose()

    if uniform_sample<prior[0]:
        '''sample from class class 1'''
        sample=A1.dot(sample_from_zero_mean_identity_covariance)+b1
        samples_class1.append(sample)
    elif (prior[0]<uniform_sample<prior[0]+prior[1]):
        '''sample from class class 2'''
        sample=A2.dot(sample_from_zero_mean_identity_covariance)+b2
        samples_class2.append(sample)
    elif (prior[0]+prior[1]<uniform_sample<prior[0]+prior[1]+prior[2
]):
        '''sample from class class 3'''
        sample=A3.dot(sample_from_zero_mean_identity_covariance)+b3
        samples_class3.append(sample)
    else :
        '''sample from class class 4'''
        sample=A4.dot(sample_from_zero_mean_identity_covariance)+b4
        samples_class4.append(sample)

samples_class1_final=np.hstack(samples_class1)
samples_class2_final=np.hstack(samples_class2)
samples_class3_final=np.hstack(samples_class3)
samples_class4_final=np.hstack(samples_class4)

a=np.squeeze(np.asarray(samples_class1_final.transpose()[0,:]))
b=np.squeeze(np.asarray(samples_class1_final.transpose()[1,:]))

c=np.squeeze(np.asarray(samples_class2_final.transpose()[0,:]))
d=np.squeeze(np.asarray(samples_class2_final.transpose()[1,:]))

e=np.squeeze(np.asarray(samples_class3_final.transpose()[0,:]))
f=np.squeeze(np.asarray(samples_class3_final.transpose()[1,:]))

g=np.squeeze(np.asarray(samples_class4_final.transpose()[0,:]))
h=np.squeeze(np.asarray(samples_class4_final.transpose()[1,:]))
plt.xlabel('Variable x1',size=13)
plt.ylabel('Variable x2',size=13)
fig.suptitle('Plot of data sampled from 4 gaussians', fontsize=15)
plt.subplots_adjust(top=.95)
plt.scatter(b,a,color='r',marker='*',label='class 1',s=50)
plt.scatter(d,c,color='g',marker='*',label='class 2',s=50)
plt.scatter(f,e,color='b',marker='*',label='class 3',s=50)
plt.scatter(h,g,color='y',marker='*',label='class 4',s=50)
plt.legend(loc='best')
plt.show()

```

```
data=np.hstack([samples_class1_final,samples_class2_final,samples_class3_final,samples_class4_final])  
'''  
data.shape=(2,samples)  
'''  
return data
```

```

In [4]: def gmm_em_kfold(data):
        '''
        create 8 models of GMM
        '''
        n_components = np.arange(1, 9)
        models = [mixture.GaussianMixture(n, covariance_type='full', random_
state=0)
                    for n in n_components]
        '''
        Lists log_likelihood,bic,aic to store result and plot
        '''
        log_likelihood=[]
        bic=[]
        aic=[]
        '''
        KFold function from sklearn
        '''
        cv = KFold(n_splits=10, random_state=42, shuffle=True)
        for m in models:
            '''
            Lists scores,scores_bic,scores_aic to store 'k' scores on 'k' di
fferent validation set
            '''
            scores = []
            scores_bic=[]
            scores_aic=[]
            for train_index, test_index in cv.split(data.T):
                X_train, X_test= data.T[train_index], data.T[test_index]
                m.fit(X_train) # fit the model
                scores.append(compute_log_likelihood(m,X_test)) # find log_
likelihood on test set and add it to the list
                scores_bic.append(m.bic(X_test)) # find bic on test set and
add it to the list
                scores_aic.append(m.aic(X_test)) # find aic on test set and
add it to the list
                log_likelihood.append(-1*sum(scores)/len(scores)) # average the
score over k validation sets
                bic.append(sum(scores_bic)/len(scores_bic)) # average the score
over k validation sets
                aic.append(sum(scores_aic)/len(scores_aic)) # average the score
over k validation sets
            from matplotlib.pyplot import figure
            fig = plt.figure(figsize=(10,10));
            plt.plot(n_components, log_likelihood, label='negative log_likelihoo
d')
            plt.xlabel('n_components',size=13);
            plt.ylabel('score',size=13);
            plt.legend(loc='best')
            fig.suptitle('Negative log_likelihood score for different GMM model
s', fontsize=15)
            plt.subplots_adjust(top=.95)
            plt.show()
            from matplotlib.pyplot import figure
            fig = plt.figure(figsize=(10,10));
            plt.xlabel('n_components',size=13);
            plt.ylabel('score',size=13);

```

```
plt.plot(n_components, bic, label='BIC')
plt.plot(n_components, aic, label='AIC')
fig.suptitle('BIC and AIC scores for different GMM models', fontsize
=15)
plt.subplots_adjust(top=.95)
plt.legend(loc='best')
plt.show()
return
```

Answer part 1. True GMM

```
In [5]: prior=[0.19,0.21,0.36,0.24]
# prior of Gaussian 1 = prior[0]
# prior of Gaussian 2 = prior[1]
# prior of Gaussian 3 = prior[2]
# prior of Gaussian 4 = prior[3]

# mean and covariance of gaussian 1
sig1=[[.15,-.1],[.1,.15]]
u1=[-1,-1]


# mean and covariance of gaussian 2
sig2=[[.15,.1],[.1,.15]]
u2=[1,1]


# mean and covariance of gaussian 3
sig3=[[.15,.1],[-.1,.15]]
u3=[-1,1]


# mean and covariance of gaussian 4
sig4=[[.15,.1],[-.1,.15]]
u4=[1,-1]
```

```
In [6]: for i in [10,100, 1000, 10000]:  
        print ("Number of samples: ",i)  
        data=generate_GMM_samples(prior=prior,number_of_samples=i,sig1=sig1,  
        \  
        sig2=sig2,sig3=sig3,sig4=sig4,u1=u1,u2=u2,u3=u3,u4=  
        u4);  
        gmm_em_kfold(data)
```


Number of samples: 10










Number of samples: 100










Number of samples: 1000









Number of samples: 10000







Conclusion

With number of samples=1000 or 10000, log likelihood and bic/aic indicate number of components = 4 is a good model for this data. Further, at $n_components=4$, the score quickly falls and plateaus around 4. Considering Occam's razor, number of components = 4 is clearly the best option

Answer 2

```
In [7]: from scipy import linalg
import numpy as np
import matplotlib.pyplot as plt
import matplotlib as mpl
from matplotlib import colors

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
In [8]: cmap = colors.LinearSegmentedColormap(
    'red_blue_classes',
    {'red': [(0, 1, 1), (1, 0.7, 0.7)],
     'green': [(0, 0.7, 0.7), (1, 0.7, 0.7)],
     'blue': [(0, 0.7, 0.7), (1, 1, 1)]})
plt.cm.register_cmap(cmap=cmap)
```

```

In [9]: def generate_GMM_samples_2class(prior,number_of_samples,sig1,sig2,u1,u2
):
    '''
    Args:
    prior of class 1 = prior[0]
    prior of class 2 = 1-prior[0]

    number_of_samples

    class 1- u_1, sig_1
    class 2- u_2, sig_2

    x is samples from zero-mean identity-covariance Gaussian sample gene
rators

    generating class 1- A1*x+b1
    generating class 2- A2*x+b2

    '''
    from matplotlib.pyplot import figure
    txt="Plot of data sampled from 4 gaussians "
    fig = plt.figure(figsize=(15,10));
    #fig.text(.35,0.09,txt,fontsize=15);

    samples_class1=[]
    samples_class2=[]

    sig_1=np.matrix(sig1)
    sig_2=np.matrix(sig2)

    u_1=np.matrix(u1).transpose()
    u_2=np.matrix(u2).transpose()

    prior=prior
    A1=np.linalg.cholesky(sig_1)
    b1=u_1

    A2=np.linalg.cholesky(sig_2)
    b2=u_2

    zero_mean=[0,0]
    cov=[[1,0],[0,1]]

    for i in range(number_of_samples):
        uniform_sample=np.random.uniform()

        sample_from_zero_mean_identity_covariance=np.random.multivariate
_normal(zero_mean,cov,[1]).transpose()

        if uniform_sample<prior[0]:
            '''sample from class class 1'''

```

```

        sample=A1.dot(sample_from_zero_mean_identity_covariance)+b1
        samples_class1.append(sample)
    elif (prior[0]<uniform_sample<prior[0]+prior[1]):
        '''sample from class class 2'''
        sample=A2.dot(sample_from_zero_mean_identity_covariance)+b2
        samples_class2.append(sample)

samples_class1_final=np.hstack(samples_class1)
samples_class2_final=np.hstack(samples_class2)

a=np.squeeze(np.asarray(samples_class1_final.transpose()[:,1]))
b=np.squeeze(np.asarray(samples_class1_final.transpose()[:,0]))

c=np.squeeze(np.asarray(samples_class2_final.transpose()[:,1]))
d=np.squeeze(np.asarray(samples_class2_final.transpose()[:,0]))
plt.xlabel('Variable x1',size=13)
plt.ylabel('Variable x2',size=13)
fig.suptitle('Data from 2 Gaussians', fontsize=15)
plt.scatter(b,a,color='r',marker='*',label='class 1',s=50)
plt.scatter(d,c,color='g',marker='*',label='class 2',s=50)
X=np.hstack([samples_class1_final,samples_class2_final])
plt.subplots_adjust(top=.95)
y = np.hstack((np.zeros(samples_class1_final.shape[1]), np.ones(samples_class2_final.shape[1])))
X=X.T
X=np.squeeze(np.asarray(X))
return X,y

```

True GMM

```

In [10]: # prior of Gaussian 1 = prior[0]
# prior of Gaussian 2 = prior[1]= 1-prior[0]
prior=[0.7,0.3]

# mean and covariance of gaussian 1
sig1=[[.15,.1],[-.1,.15]]
u1=[2,-1]

# mean and covariance of gaussian 2
sig2=[[.15,.1],[.1,1.5]]
u2=[1.2,-1]
X,y=generate_GMM_samples_2class(prior=prior,number_of_samples=999,sig1=sig1,sig2=sig2,\
                                u1=u1,u2=u2);


print ("Samples from GMM :",X.shape)
print ("Class labels of the GMM :",y.shape)

```

```

Samples from GMM : (999, 2)
Class labels of the GMM : (999,)

```



LDA

```

In [11]: def plot_data(lda, X, y, y_pred):
'''
    LDA helper function
'''
plt.xlabel('Variable x1',size=13)
plt.ylabel('Variable x2',size=13)
fig.suptitle('LDA decision boundary for data from 2 Gaussians', font
size=15)
tp = (y == y_pred) # True Positive
tp0, tp1 = tp[y == 0], tp[y == 1] # True Positive for class 0 and cl
ass 1 respectively
# tp0 is boolean with "tp and class=0"
X0, X1 = X[y == 0], X[y == 1] # points
X0_tp, X0_fp = X0[tp0], X0[~tp0]
X1_tp, X1_fp = X1[tp1], X1[~tp1]

# class 0: dots
plt.scatter(X0_tp[:, 0], X0_tp[:, 1], marker='.', color='red')
plt.scatter(X0_fp[:, 0], X0_fp[:, 1], marker='x',
            s=20, color='#990000') # dark red

# class 1: dots
plt.scatter(X1_tp[:, 0], X1_tp[:, 1], marker='.', color='blue')
plt.scatter(X1_fp[:, 0], X1_fp[:, 1], marker='x',
            s=20, color='#000099') # dark blue

# class 0 and 1 : areas
nx, ny = 200, 100
x_min, x_max = plt.xlim()
y_min, y_max = plt.ylim()
xx, yy = np.meshgrid(np.linspace(x_min, x_max, nx),
                     np.linspace(y_min, y_max, ny))
Z = lda.predict_proba(np.c_[xx.ravel(), yy.ravel()])
Z = Z[:, 1].reshape(xx.shape)
plt.pcolormesh(xx, yy, Z, cmap='red_blue_classes',
               norm=colors.Normalize(0., 1.), zorder=0)
plt.contour(xx, yy, Z, [0.5], linewidths=2., colors='white')

# means
plt.plot(lda.means_[0][0], lda.means_[0][1],
        '*', color='yellow', markersize=15, markeredgcolor='grey')
plt.plot(lda.means_[1][0], lda.means_[1][1],
        '*', color='yellow', markersize=15, markeredgcolor='grey')

return

```

```

In [12]: lda = LinearDiscriminantAnalysis(store_covariance=True)
lda.fit(X, y)
y_pred = lda.predict(X)
print ("Accuracy with LDA is :",np.mean(y_pred==y))

```

Accuracy with LDA is : 0.8858858858858859


```

In [13]: '''
> To Predict-
>> ((np.matmul(X,lda.coef_.T)+ lda.intercept_)>0).astype(int)
>>or
>> lda.predict(X).reshape(-1,1);
'''

# LDA coefficients are-
w=lda.coef_.T
b=lda.intercept_
print (w,b)


[[-5.531773]
 [-0.419741]] [7.53545276]

```

```


In [14]: from matplotlib.pyplot import figure
fig = plt.figure(figsize=(15,10));
fig.suptitle('Linear Discriminant Analysis', fontsize=15)
plot_data(lda, X, y, y_pred)
plt.subplots_adjust(top=.95)
plt.show()

```



Logistic Regression

```
In [15]: def test_logistic_function(x):  
         return 1/(1+np.exp(x))  
x=np.arange(-10,10,.1)  
plt.plot(x,test_logistic_function(x));
```



Answer 2

$$\rightarrow W_{LDA} \cdot X + b_{LDA} \sum_{\substack{\text{class 1} \\ \text{class 0}}} 0$$

$$\gamma(x) = \frac{1}{1 + e^{-(w^T x + b)}}$$

$$\begin{pmatrix} W_{MLE} \\ b_{MLE} \end{pmatrix} = \underset{\begin{pmatrix} W \\ b \end{pmatrix}}{\operatorname{argmax}} P[D | w]$$

where $D = \begin{bmatrix} x_1, L_1 \\ x_2, L_2 \\ \vdots \\ x_N, L_N \end{bmatrix}$; $x_i \in \mathbb{R}^2 \rightarrow \text{gram sample}$
 $L_i \in \{0, 1\}$

as samples are i.i.d \rightarrow

$$\begin{pmatrix} W_{MLE} \\ b_{MLE} \end{pmatrix} = \underset{\begin{pmatrix} W \\ b \end{pmatrix}}{\operatorname{argmax}} \sum_{i=1}^N \log P \left(\begin{matrix} x_i \\ L_i \end{matrix} \middle| \begin{matrix} W \\ b \end{matrix} \right)$$

$$\begin{pmatrix} W_{MLE} \\ b_{MLE} \end{pmatrix} = \underset{\begin{pmatrix} W \\ b \end{pmatrix}}{\operatorname{argmax}} \left[\sum_{i=1}^N \log p(L_i | x_i, W, b) + \sum_{i=1}^N \log p(x_i | W, b) \right]$$

$$p(x_i | W, b) = p(x_i) \quad (\text{independent of } W, b)$$

$$\Rightarrow \begin{pmatrix} W_{MLE} \\ b_{MLE} \end{pmatrix} = \underset{\begin{pmatrix} W \\ b \end{pmatrix}}{\operatorname{argmax}} \sum_{i=1}^N \log p(L_i | x_i, W, b)$$

~~$p(L_i | x_i, W, b)$~~ Assuming class labels
0 and 1 \rightarrow

$$\Rightarrow \begin{pmatrix} W_{MLE} \\ b_{MLE} \end{pmatrix} = \underset{\begin{pmatrix} W \\ b \end{pmatrix}}{\operatorname{argmax}} \sum_{i=1}^N \log \left[y(x_i)^{L_i} (1 - y(x_i))^{(1-L_i)} \right]$$

$$= \underset{\begin{pmatrix} W \\ b \end{pmatrix}}{\operatorname{argmax}} \sum_{i=1}^N \left[L_i \log y(x_i) + (1-L_i) \log (1 - y(x_i)) \right]$$

Let $\gamma(x) = \frac{1}{1 + e^x} \rightarrow$ sigmoid fⁿ for vs.

$$\gamma'(x) = \frac{-1}{(1 + e^x)^2} \times e^x$$

$$\begin{aligned} \gamma'(x) &= \gamma(x) \left[\frac{1}{1 + e^x} - 1 \right] \\ &= \gamma(x) [\gamma(x) - 1] \end{aligned}$$

$$\gamma'(x) = -\gamma(x)(1 - \gamma(x))$$

further represent $\begin{bmatrix} w \\ b \end{bmatrix}$ by a single parameter θ

and add scalar '1' to the feature 'x';

so that $\left[w^T x + b = \theta^T x \right]$

$$\Rightarrow \theta_{MLE} = \underset{\theta}{\operatorname{argmax}} \sum_{i=1}^N \left[L_i \log \gamma(\theta^T x) + (1 - L_i) \log (1 - \gamma(\theta^T x)) \right]$$

$L(\theta) \leftarrow$

To maximize ; use gradient ascent \rightarrow

$$\theta = \theta + \alpha \nabla_{\theta} J(\theta)$$

for 1 training example \rightarrow

$$\frac{\partial J(\theta)}{\partial \theta_j} = \left[\frac{L_i}{y(\theta^T x)} - \frac{(1-L_i)}{[1-y(\theta^T x)]} \right] \frac{\partial g(\theta^T x)}{\partial \theta_j}$$

where $j = 1, 2, \dots$

$$\theta = [w_1, w_2, b]$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = - \left[\frac{L_i}{y(\theta^T x)} - \frac{(1-L_i)}{(1-y(\theta^T x))} \right] g(\theta^T x) \times$$

$$g(1-g(\theta^T x)) \times \frac{\partial \theta^T x}{\partial \theta_j}$$

$$= - \left[L_i(1-y(\theta^T x)) - (1-L_i)y(\theta^T x) \right] x_{ij}$$

$$= - [L_i - y(\theta^T x)] x_{ij} \quad ; \quad \begin{array}{l} i \rightarrow \text{example} \\ j \rightarrow \text{feature number} \end{array}$$

$$\Rightarrow \theta_j = \theta_j - 2 (y^i - y(\theta^T x^i)) x_j^i$$

```

In [16]: class LogisticRegression:
    def __init__(self, lr=0.01, num_iter=100000, fit_intercept=True, verbose=True):
        self.lr = lr
        self.num_iter = num_iter
        self.fit_intercept = fit_intercept
        self.verbose = verbose

    def __add_intercept(self, X):
        intercept = np.ones((X.shape[0], 1))
        return np.concatenate((intercept, X), axis=1)

    def __sigmoid(self, z):
        return 1 / (1 + np.exp(-z))

    def __loss(self, h, y):
        return (y * np.log(h) + (1 - y) * np.log(1 - h)).mean()

    def fit(self, X, y):
        if self.fit_intercept:
            X = self.__add_intercept(X)

        # weights initialization
        self.theta = np.zeros(X.shape[1])

        for i in range(self.num_iter):
            z = np.dot(X, self.theta)
            h = self.__sigmoid(z)
            gradient = np.dot(X.T, (y - h)) / y.size
            self.theta = self.theta - self.lr * gradient

            z = np.dot(X, self.theta)
            h = self.__sigmoid(z)
            loss = self.__loss(h, y)

            if (self.verbose == True and i % 10000 == 0):
                print(f'loss: {loss} \t')

    def predict_prob(self, X):
        if self.fit_intercept:
            X = self.__add_intercept(X)

        return self.__sigmoid(np.dot(X, self.theta))

    def predict(self, X):
        return self.predict_prob(X).round()

```



```

In [17]: def plot_logistic():
    '''
    Helper function
    '''

    from matplotlib.pyplot import figure
    fig = plt.figure(figsize=(15,10));
    plt.xlabel('Variable x1',size=13)
    plt.ylabel('Variable x2',size=13)
    fig.suptitle('Logistic Regression decision boundary for data from 2
Gaussians', fontsize=15)
    tp = (y == y_pred) # True Positive
    tp0, tp1 = tp[y == 0], tp[y == 1] # True Positive for class 0 and class 1 respectively
    # tp0 is boolean with "tp and class=0"
    X0, X1 = X[y == 0], X[y == 1] # points
    X0_tp, X0_fp = X0[tp0], X0[~tp0]
    X1_tp, X1_fp = X1[tp1], X1[~tp1]

    # class 0: dots
    plt.scatter(X0_tp[:, 0], X0_tp[:, 1], marker='.', color='red')
    plt.scatter(X0_fp[:, 0], X0_fp[:, 1], marker='x',
                s=20, color='#990000') # dark red

    # class 1: dots
    plt.scatter(X1_tp[:, 0], X1_tp[:, 1], marker='.', color='blue')
    plt.scatter(X1_fp[:, 0], X1_fp[:, 1], marker='x',
                s=20, color='#000099') # dark blue

    # class 0 and 1 : areas
    nx, ny = 200, 100
    x_min, x_max = plt.xlim()
    y_min, y_max = plt.ylim()
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, nx),
                          np.linspace(y_min, y_max, ny))
    Z=model.predict_prob((np.c_[xx.ravel(), yy.ravel()])))
    Z=Z.reshape(xx.shape)

    plt.pcolormesh(xx, yy, Z, cmap='red_blue_classes',
                   norm=colors.Normalize(0., 1.), zorder=0)
    plt.contour(xx, yy, Z, [0.5], linewidths=2., colors='white');
    plt.subplots_adjust(top=.95);

```

```

In [18]: model = LogisticRegression(lr=0.0042, num_iter=int(42000*4.2))
X=np.squeeze(np.asarray(X))

```

```
In [19]: model.fit(X, y)
```

```
loss: -0.6913611622242243
loss: -0.3925256110835711
loss: -0.34937699814411055
loss: -0.3271342605257729
loss: -0.31430279830677393
loss: -0.30628444660081783
loss: -0.30098068931434596
loss: -0.29732343894435354
loss: -0.29472103651588316
loss: -0.29282351110119864
loss: -0.2914128817526874
loss: -0.2903476447177061
loss: -0.2895327958119807
loss: -0.2889027510410482
loss: -0.2884111704433018
loss: -0.2880246612358051
loss: -0.28771875344985876
loss: -0.2874752552538597
```

```
In [20]: y_pred = model.predict(X)
print ("Accuracy of logistic regression is :", (y_pred == y).mean())
```

```
Accuracy of logistic regression is : 0.8828828828828829
```

```
In [21]: plot_logistic()
```



MAP

Maximize posterior with true priors and probabilities. Select the class with max posterior

```
In [22]: '''
Helper function
'''

def plot_map(X,y,y_pred):
    from matplotlib.pyplot import figure
    fig = plt.figure(figsize=(15,10));
    plt.xlabel('Variable x1',size=13)
    plt.ylabel('Variable x2',size=13)
    fig.suptitle('MAP decision boundary for data from 2 Gaussians', font
size=15)
    tp = (y == y_pred) # True Positive
    tp0, tp1 = tp[y == 0], tp[y == 1] # True Positive for class 0 and cl
ass 1 respectively
    # tp0 is boolean with "tp and class=0"
    X0, X1 = X[y == 0], X[y == 1] # points
    X0_tp, X0_fp = X0[tp0], X0[~tp0]
    X1_tp, X1_fp = X1[tp1], X1[~tp1]

    # class 0: dots
    plt.scatter(X0_tp[:, 0], X0_tp[:, 1], marker='.', color='red')
    plt.scatter(X0_fp[:, 0], X0_fp[:, 1], marker='x',
                s=20, color='#990000') # dark red

    # class 1: dots
    plt.scatter(X1_tp[:, 0], X1_tp[:, 1], marker='.', color='blue')
    plt.scatter(X1_fp[:, 0], X1_fp[:, 1], marker='x',
                s=20, color='#000099') # dark blue

    # class 0 and 1 : areas
    nx, ny = 200, 100
    x_min, x_max = plt.xlim()
    y_min, y_max = plt.ylim()
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, nx),
                        np.linspace(y_min, y_max, ny))

    Z = (gaussian1.pdf((np.c_[xx.ravel(), yy.ravel()])*prior[1])/((gaus
sian2.pdf((np.c_[xx.ravel(), yy.ravel()])*prior[1])+(gaussian1.pdf((np.
c_[xx.ravel(), yy.ravel()])*prior[1]))
    Z=Z.reshape(xx.shape)

    plt.pcolormesh(xx, yy, Z, cmap='red_blue_classes',
                    norm=colors.Normalize(0., 1.), zorder=0)
    plt.contour(xx, yy, Z, [0.5], linewidths=2., colors='white');
    plt.subplots_adjust(top=.95);
```

```

In [23]: # Re writing true priors-

# prior of Gaussian 1 = prior[0]
# prior of Gaussian 2 = prior[1]= 1-prior[0]
prior=[0.7,0.3]

# mean and covariance of gaussian 1
sig1=[[.15,.1],[-.1,.15]]
u1=[2,-1]

# mean and covariance of gaussian 2
sig2=[[.15,.1],[.1,1.5]]
u2=[1,-1]

print ("Samples from GMM :",X.shape)
print ("Class labels of the GMM :",y.shape)

```

```

Samples from GMM : (999, 2)
Class labels of the GMM : (999,)

```

```

In [24]: gaussian1 = multivariate_normal(mean=u1, cov=sig1)
gaussian2 = multivariate_normal(mean=u2, cov=sig2)

class1_posterior=(gaussian1.pdf(X)*prior[0])/((gaussian1.pdf(X)*prior[0]
)+(gaussian2.pdf(X)*prior[0]))
class2_posterior=(gaussian2.pdf(X)*prior[0])/((gaussian1.pdf(X)*prior[0]
)+(gaussian2.pdf(X)*prior[0]))

class1_class2_posterior_stacked=np.vstack([class1_posterior,class2_poste
rior]).T

y_pred=np.argmax(class1_class2_posterior_stacked,axis=1)
print ("Accuracy of MAP is :",np.mean(y==y_pred))


```

```

Accuracy of MAP is : 0.9419419419419419

```

```
In [25]: plot_map(X,y,y_pred)
```



Conclusion

MAP classifier has the highest accuracy of close to 94%. LDA and logistic regression have accuracy close to 88%. The decision boundaries of 3 classifiers are shown above

References

1. [Lecture notes \(https://github.com/rohinarora/EECE5644-Machine_Learning\)](https://github.com/rohinarora/EECE5644-Machine_Learning)
2. https://scikit-learn.org/0.16/auto_examples/classification/plot_lda_qda.html (https://scikit-learn.org/0.16/auto_examples/classification/plot_lda_qda.html)
3. https://github.com/martinpella/logistic-reg/blob/master/logistic_reg.ipynb (https://github.com/martinpella/logistic-reg/blob/master/logistic_reg.ipynb)
4. <http://cs229.stanford.edu/notes/cs229-notes1.pdf> (<http://cs229.stanford.edu/notes/cs229-notes1.pdf>)
5. [MinM \(https://en.wikipedia.org/wiki/Eminem\)](https://en.wikipedia.org/wiki/Eminem)
6. [Don't Panic \(https://en.wikipedia.org/wiki/Don't_Panic:_The_Official_Hitchhiker's_Guide_to_the_Galaxy_Companion\)](https://en.wikipedia.org/wiki/Don't_Panic:_The_Official_Hitchhiker's_Guide_to_the_Galaxy_Companion)